





Incidents and Chatops

Manage your incidents with Stride & Hipchat



RAMIRO BERRELLEZA | ARCHITECT | @RBERRELLEZA

Agenda

Incidents

Chatops

Before the incident

During the incident

After the incidents

Resources available

**Incident: An interruption in
your service that requires
human intervention**




Availability % ⇅	Downtime per year ⇅	Downtime per month ⇅	Downtime per week ⇅	Downtime per day ⇅
90% ("one nine")	36.5 days	72 hours	16.8 hours	2.4 hours
95% ("one and a half nines")	18.25 days	36 hours	8.4 hours	1.2 hours
97%	10.96 days	21.6 hours	5.04 hours	43.2 minutes
98%	7.30 days	14.4 hours	3.36 hours	28.8 minutes
99% ("two nines")	3.65 days	7.20 hours	1.68 hours	14.4 minutes
99.5% ("two and a half nines")	1.83 days	3.60 hours	50.4 minutes	7.2 minutes
99.8%	17.52 hours	86.23 minutes	20.16 minutes	2.88 minutes
99.9% ("three nines")	8.76 hours	43.8 minutes	10.1 minutes	1.44 minutes
99.95% ("three and a half nines")	4.38 hours	21.56 minutes	5.04 minutes	43.2 seconds
99.99% ("four nines")	52.56 minutes	4.38 minutes	1.01 minutes	8.64 seconds
99.995% ("four and a half nines")	26.28 minutes	2.16 minutes	30.24 seconds	4.32 seconds
99.999% ("five nines")	5.26 minutes	25.9 seconds	6.05 seconds	864.3 milliseconds
99.9999% ("six nines")	31.5 seconds	2.59 seconds	604.8 milliseconds	86.4 milliseconds
99.99999% ("seven nines")	3.15 seconds	262.97 milliseconds	60.48 milliseconds	8.64 milliseconds
99.999999% ("eight nines")	315.569 milliseconds	26.297 milliseconds	6.048 milliseconds	0.864 milliseconds
99.9999999% ("nine nines")	31.5569 milliseconds	2.6297 milliseconds	0.6048 milliseconds	0.0864 milliseconds

**Incidents are (mostly)
unavoidable**



**What really matters is how
we react when they do
happen**



Agenda

Incidents

Chatops

Before the incident

During the incident

After the incidents

Resources available

ChatOps is a collaboration model that connects people, tools, process, and automation into a transparent workflow.

**It connects teams and all the
systems you rely on in a
single location**





For all teams



**For teams
behind the
firewall**

Agenda

Incidents

Chatops

Before the incident

During the incident

After the incidents

Resources available



Have a plan

Checklist

List all the actions that you want accomplished during an incident. Be as detailed as you can.

Ownership

Every action in your plan should be owned by a single person

Train

Simulate an incident a few times, so you polish your plan

Have a plan

Checklist

List all the actions that you want accomplished during an incident. Be as detailed as you can

Ownership

Every action in your plan should be owned by a single person

Train

Simulate an incident a few times, so you polish your plan

Have a plan

Checklist

List all the actions that you want accomplished during an incident. Be as detailed as you can

Ownership

Every action in your plan should be owned by a single person

Train

Simulate an incident a few times, so you polish your plan

automate



all the things!

Automate

Health checks

All your services should report their health frequently

Monitor

Use a system to keep an eye on your health checks. Also, setup some dashboards

Notify

When health checks fail, automatically notify your systems

Automate

Health checks

All your services should report their health frequently

Monitor

Use a system to keep an eye on your health checks. Also, setup some dashboards

Notify





When health checks fail, automatically notify your systems

pagerduty



DATADOG

Nagios®





Pro tip: Don't run your monitoring infra on the same infra as the system you're monitoring

Automate

Health checks

All your services should report their health frequently

Monitor

Use a system to keep an eye on your health checks. Also, setup some dashboards

Notify

When health checks fail, automatically notify your systems





Statuspage



Datadog 3:56 PM

! [No data on {host:db-208.sc1.bb-inf.net,process:postgres}] [bb-cloud] PostgreSQL is not running ...

There is currently no data being reported for the PostgreSQL process on [db-208.sc1.bb-inf.net](#). @stride-hello-bitbucket.cloud.alerts

host:db-208.sc1.bb-inf.net,process:postgres has been missing data for the last 10m

! [Triggered on {host:db-208.sc1.bb-inf.net}] [bb-cloud] Datadog agent not reporting

[db-208.sc1.bb-inf.net](#) / with IP 10.193.8.64 is not reporting
@stride-hello-bitbucket.cloud.alerts"



PagerDuty 4:06 PM



PagerDuty / Bitbucket Datadog

Triggered #1678568

No Data: [bb-cloud] PostgreSQL Replication Lag for SC1 Disaster Recovery Databases...

Respond

Assigned: Ching Sun Urgency:  ↑ High

Agenda

Incidents

Chatops


Before the incident

During the incident

After the incidents

Resources available





Create a single point of truth

Canonical information

Everybody involved will refer back to this. Link to all the other resources here.



Keep it simple

Only use it for very significant updates.

Keep it updated

Ensure that this always displays the latest state of your incident

Create a single point of truth

Canonical information

Everybody involved will refer back to this. Link to all the other resources here.

Keep it simple

Only use it for very significant updates.

Keep it updated

Ensure that this always displays the latest state of your incident

Create a single point of truth

Canonical information

Everybody involved will refer back to this. Link to all the other resources here.

Keep it simple

Only use it for very significant updates.

Keep it updated

Ensure that this always displays the latest state of your incident

Better comms

Cannonical point of contact

All the communication should be reflected on the room

Keep it private


Only invite people that are actively involved in the incident

Over communicate

Share all the information, tasks, questions and answers

 Coconut

ROOMS

 INCIDENT-3541 HC Data Center Eng HipChat Server/DC Supp... HC Data Center Social HC Data Center Crowd HCDC PM Dev (Hip)Chat...

INCIDENT-3541

<https://confluence.atlassian.com/incidents/3541> | POC rberrelleza@atlassian.com

Ramiro Berrelleza · Aug-5 1:47 PM

**Click here to join the video call**hipchat.me

Ramiro Berrelleza has started a room video call. Click the link above to join the call.

Friday August 11, 2017



Ramiro Berrelleza · Aug-11 4:24 PM

patch.tar.gz 14K

Better comms

Cannonical point of contact

All the communication should be reflected on the room

Keep it private

Only invite people that are actively involved in the incident



Over communicate

Share all the information, tasks, questions and answers

Better comms

Cannonical point of contact

All the communication should be reflected on the room

Keep it private

Only invite people that are actively involved in the incident

Over communicate

Share all the information, tasks, questions and answers

A screenshot of a mobile application interface. A circular callout highlights a specific item in a list. The item has a blue icon with a white camera-like symbol and a star icon. The text "ing now" is visible above the item, and "5-5295-4b87..." is visible below it. The background is a light blue gradient.



Pro tip: Let everybody know when the incident is over.

Archive your rooms if possible for future reference.

Agenda

Incidents

Chatops

Before the incident

During the incident

After the incidents

Resources available

Postmortem

Root case analysis

What really happened?

Improvements

Create and assign tickets with concrete tasks that will prevent this from happening again

Communication

Build trust with your users by communicating the results of the incidents

Postmortem

Root case analysis

What really happened?

Improvements

Create and assign tickets with concrete tasks that will prevent this from happening again

Communication

Build trust with your users by communicating the results of the incidents

Postmortem

Root case analysis

What really happened?

Improvements

Create and assign tickets with concrete tasks that will prevent this from happening again

Communication

Build trust with your users by communicating the results of the incidents, both internally and externally

February 2018

US region network outage

We experienced a network outage from 03:48 to 03:54 UTC and 04:03 to 04:10 UTC in our US region which impacted a ...
Feb 27, 05:13 - 05:13 UTC

Jira Kanban Board Issues

Some Jira customers experienced problems viewing Issues from Kanban Boards between 23:05 UTC Feb 21 2018 and ...
Feb 22, 07:38 - 07:38 UTC

Degraded Services for Confluence Cloud

This issue has been resolved. Confluence Cloud services are fully operational.
Feb 21, 15:00 - 22:09 UTC

Confluence Questions

The engineering team has declared the incident over.
Feb 21, 10:18 - 13:55 UTC

Bitbucket Issue Create and Edit Screens

This incident has been resolved.
Feb 21, 06:41 - 08:05 UTC

General outage in Atlassian Cloud

The incident is now resolved as systems are operational and all monitoring is back to normal.
Feb 20, 23:00 - 23:23 UTC

Degraded service to Jira Cloud in EU and US Regions

The issue impacting Jira and Confluence Cloud in EU has been resolved and all services have been verified to be opera...
Feb 16, 13:03 - 20:01 UTC

Jira Cloud outage in US East Region

This incident has been resolved.
Feb 7, 00:11 - 05:18 UTC

Agenda

Incidents

Chatops

Before the incident

During the incident

After the incidents

Resources available

Playbooks

Incident response communications

<https://www.atlassian.com/team-playbook/plays/incident-response-communications>

Pre mortem

<https://www.atlassian.com/team-playbook/plays/pre-mortem>

5-whys analysis

<https://www.atlassian.com/team-playbook/plays/5-whys>

Stride integrations

Use status page and stride together

<https://confluence.atlassian.com/stride-documentation/use-statuspage-and-stride-together-945522458.html>

Pagerduty's Stride extension guide

<https://www.pagerduty.com/docs/guides/stride-extension-guide/>

Make your own bot and integrations

Stride developer API

<https://developer.atlassian.com/cloud/stride/>

Convert your Hipchat addons to Stride apps

<https://developer.atlassian.com/cloud/stride/learning/hipchat-to-stride/>

Sample stride bot in glitch

<https://glitch.com/~stride-glitch>

Inside Atlassian

Inside Atlassian: how IT & SRE use ChatOps to run incident management

<https://www.atlassian.com/blog/it-service-management/inside-atlassian-sre-use-chatops-run-incident-management>

In summary ...



Prepare

Have an incident response plan ready.



Execute

Follow your plan, solve the incident quickly; you got this.



Prevent

Understand what happened, and improve your system.

QA



Thanks



RAMIRO BERRELLEZA | ARCHITECT | @RBERRELLEZA

